

Virtual Memory

Virtual memory is a memory management technique that creates an **illusion of a large memory space** for programs, even if the actual physical memory is smaller. It allows programs to use more memory than is physically available by temporarily transferring data between **RAM** and **disk storage**.

- **Logical Address Space:** Programs operate in a logical address space, which is larger than physical memory.
- **Page Mapping:** The operating system uses a **page table** to map logical addresses to physical memory locations or to disk storage (swap space).

Advantages of Virtual Memory

1. **Efficient Memory Usage:** Programs only load the necessary parts into RAM, reducing memory wastage.
2. **Support for Larger Programs:** Enables execution of programs larger than the physical memory.
3. **Multiprogramming:** Allows multiple programs to run simultaneously, sharing physical memory.

Demand Paging

Demand paging is a lazy loading technique where program pages are loaded into memory **only when they are accessed**.

How Demand Paging Works


1. When a program requests a page, the operating system checks if it is in memory.
2. If the page is not in memory, a **page fault** occurs.
3. The operating system retrieves the page from secondary storage (disk) and loads it into physical memory.

Page Fault

A **page fault** occurs when a program tries to access a page that is not currently in physical memory.

Steps During a Page Fault

1. **Trap:** The CPU traps to the operating system.
2. **Check:** The operating system checks the page table to determine if the page is valid.
3. **Load Page:** If valid, the page is fetched from disk and loaded into memory.
4. **Update Page Table:** The page table is updated with the new page's location in memory.
5. **Resume Execution:** The process resumes from the point where the page fault occurred.



Performance of Demand Paging

The performance of demand paging depends on the **page fault rate**, which measures the frequency of page faults.

1. Page Fault Rate (p):

- **p = 0**: No page faults (ideal case).
- **p = 1**: Every access results in a page fault (worst case).

2. Effective Access Time (EAT):

The effective access time combines the cost of accessing memory and handling page faults:

$$\text{EAT} = (1-p) \times \text{Memory Access Time} + p \times (\text{Page Fault Service Time})$$

Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead} \\ &) \end{aligned}$$

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$
 $= (1 - p) \times 200 + p \times 8,000,000$
 $= 200 + p \times 7,999,800$
- If one access out of 1,000 causes a page fault, then
 $EAT = 8.2 \text{ microseconds}$.
This is a slowdown by a factor of 40!!

High page fault rates significantly degrade performance due to the slow disk access involved in handling faults.

Page Replacement

When physical memory is full and a new page needs to be loaded, the operating system must replace an existing page using a **page replacement algorithm**.

Common Page Replacement Algorithms

1. FIFO (First-In, First-Out):

- The oldest page in memory is replaced.
- Simple but may lead to **Belady's Anomaly** (more frames can increase page faults).

FIFO ALGO

F1

F2

F3

7	0	1	2	0	3	0	4	2	3	0	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2	2	2	2	4	4	4	0	0	0	0	0	0	0	0	0	7	7	7
0	0	0		3	3	3	2	2	2	2	2	2	2	1	1	1	1	1	0	0	
1	1	1	1	0	0	0	3	3	3	3	3	3	3	3	2	2	2	2	2	1	
X	X	X	X	✓	X	X	X	X	X	✓	✓	✓	✓	✓	X	X	✓	✓	X	X	X

(MISS) Fault page = 15

(HIT) Not fault page = 7

Hit ratio = $\frac{7}{22}$

Miss Ratio = $\frac{15}{22}$

2. Optimal Replacement:

- | | 7 | 0 | 1 | 2 | 0 | 3 | 0 | 4 | 2 | 3 | 0 | 3 | 2 | 1 | 2 | 0 | 1 | 7 | 0 | 1 |
|----|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| F1 | 7 | 7 | 7 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 7 | 7 | 7 |
| F2 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| F3 | | 1 | 1 | 1 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| | X | X | X | X | ✓ | X | ✓ | X | X | ✓ | X | ✓ | ✓ | X | ✓ | ✓ | ✓ | X | ✓ | ✓ |

- Replaces the page that has not been used for the longest time.
- Effective but requires additional hardware or data structures.

4